

# Reinforcement Learning

At a high level, our RL should enhance the predictive accuracy and adaptability of our stock prediction model using a reinforcement learning layer. This layer will dynamically assign weights to the outputs of our NLP, XGBoost, and LSTM models based on real-time market feedback. By adjusting the influence of each model according to market conditions and recent performance, the RL layer should optimize the combined prediction and improve profitability and accuracy.

## State

- We need to include each model's prediction for the price or probability that it increases in the state. We should also include the overall sentiment score from NLP.
- We should add multiple technical indicators and see how they impact performance because they give the RL agent context of the market environment.
  - These indicators can help show under what circumstances each model performs best and use that information to weight each model accordingly.
- We need to do some feature engineering to decide which technical indicators are most relevant.
  - For example, long-term indicators like interest rates or broad market indices that aren't explicitly used in LSTM and XGBoost could add valuable context.
- Keeping the state space simple with only the most relevant features is important to reduce overfitting.
- **Example indicators:** Volatility and a trend indicator like the slope of a moving average could provide the context needed without overwhelming the model.

Here's an example of how the state vector might be structured in code:

To implement it like this we will have to do more data processing

```
state = {
    'xgboost_prediction': xgboost_model.predict(),
    'lstm_prediction': lstm_model.predict(),
    'nlp_sentiment': nlp_model.get_sentiment_score(),
    'volatility': calculate_volatility(),
    'trend_slope': calculate_moving_average_slope(),
    'rolling_accuracy_xgboost': calculate_rolling_accuracy(xgboost_model),
    'rolling_accuracy_lstm': calculate_rolling_accuracy(lstm_model),
    'rolling_accuracy_nlp': calculate_rolling_accuracy(nlp_model)
}
```

## Action

- Weighting each model's prediction: The action space will allow the agent to assign continuous weights to each model's prediction. It should be dynamic as well.

## Reward

- The reward will be based on the profitability and or accuracy of the weighted prediction relative to the actual price movement.
- If the weighted prediction performs well, the agent should receive a positive reward to reinforce the chosen weighting strategy.

```
def calculate_reward(actual_price, combined_prediction, prev_price):
    price_change = actual_price - prev_price
    predicted_change = combined_prediction - prev_price
    direction_reward = 1 if (predicted_change * price_change > 0) else -1
    magnitude_reward = abs(predicted_change - price_change)
    return direction_reward - magnitude_reward
```

## Reinforcement Learning Algorithm

- Proximal Policy Optimization seems like the best choice based on my research. PPO is a policy gradient method and is well-suited for this application because of it's
  - Ability to handle continuous action spaces, which is ideal for weighting models with fractional values.
  - Stability, which is essential for applications in volatile markets.
  - Capability to optimize weights based on both market conditions and model performance.
  - Direct optimization of continuous values for model weights, which fits our needs perfectly.
- The Stable Baselines3 library has strong support for PPO, which should make implementation somewhat straightforward.

## Implementation Plan

### Training Phase

- We will conduct backtesting using historical data, split into training and validation sets.
- The agent will learn optimal weighting strategies based on each model's past performance and market conditions.
- Feature engineering and experimentation with technical indicators will help us refine the state representation
- <https://spinningup.openai.com/en/latest/algorithms/ppo.html>

### Evaluation Phase

- The agent's performance will be evaluated on validation data to assess its ability to generalize.
- Testing on out-of-sample data will allow us to see how well the agent is doing and ensure it can adapt to unseen market conditions